

Differentiable Programming

Dr. Svetlin Penkov

Logistic Regression


$$\mathcal{L} = - \sum_{i=1}^N y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T))$$

where

$$p(y_i = 1 | x_i, W) = \sigma(x_i W^T)$$


Computational Graph

$$\mathcal{L}_i = - \left[y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T)) \right]$$




Computational Graph

$$\mathcal{L}_i = - [y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T))]$$




Computational Graph

$$\mathcal{L}_i = - \left[y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T)) \right]$$



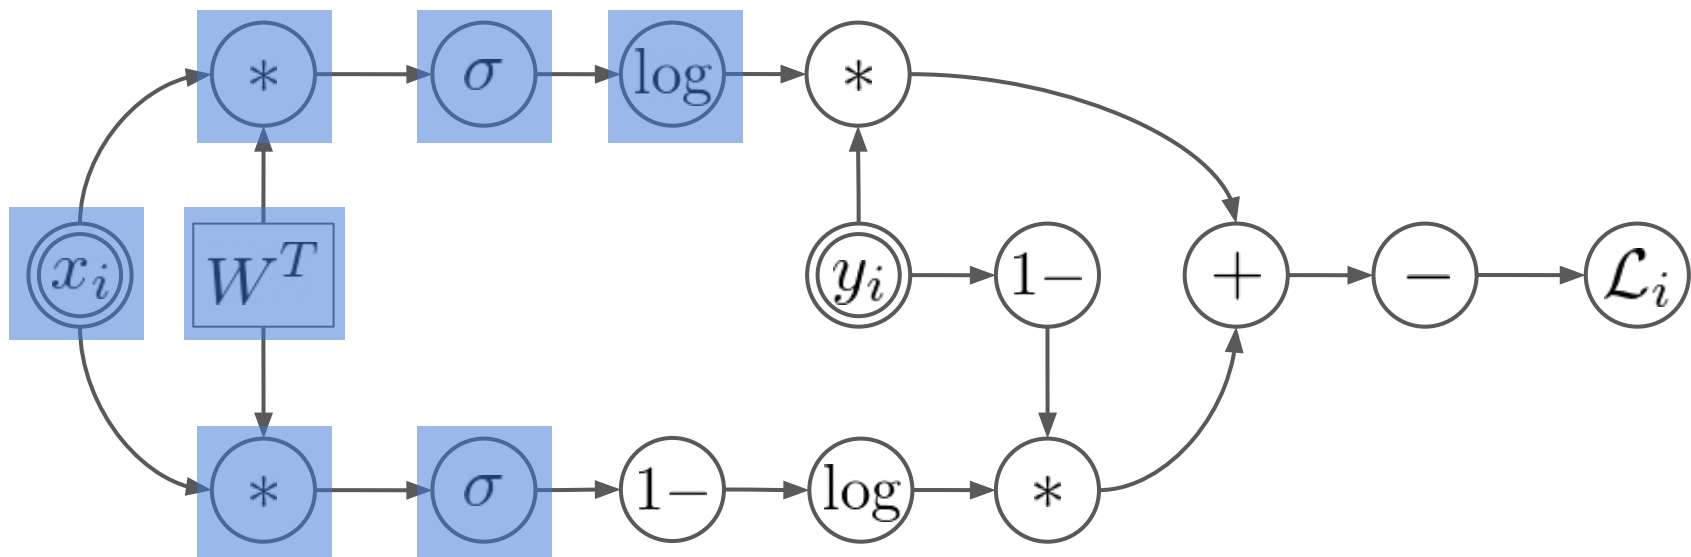
Computational Graph

$$\mathcal{L}_i = - \left[y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T)) \right]$$




Computational Graph

$$\mathcal{L}_i = - \left[y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T)) \right]$$




Computational Graph

$$\mathcal{L}_i = - \left[y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T)) \right]$$




Computational Graph

$$\mathcal{L}_i = - \left[y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T)) \right]$$




Computational Graph

$$\mathcal{L}_i = - \left[y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T)) \right]$$




Computational Graph

$$\mathcal{L}_i = - \left[y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T)) \right]$$




Computational Graph

$$\mathcal{L}_i = - [y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T))]$$




Computational Graph

$$\mathcal{L}_i = - [y_i \log \sigma(x_i W^T) + (1 - y_i) \log (1 - \sigma(x_i W^T))]$$




Computational Graph

$$\frac{\partial \mathcal{L}_i}{\partial W^T} = ?$$




Computational Graph

$$\frac{\partial \mathcal{L}_i}{\partial W^T} = ?$$




Computational Graph

$$\frac{\partial \mathcal{L}_i}{\partial W^T} = \frac{\partial \mathcal{L}_i}{\partial z} \frac{\partial z}{\partial W^T}$$




Computational Graph

$$\frac{\partial \mathcal{L}_i}{\partial W^T} = \overbrace{\frac{\partial \mathcal{L}_i}{\partial z}}^{\text{known}} \frac{\partial z}{\partial W^T}$$




Computational Graph

$$\frac{\partial \mathcal{L}_i}{\partial W^T} = \overbrace{\frac{\partial \mathcal{L}_i}{\partial z}}^{\text{known}} \overbrace{\frac{\partial z}{\partial W^T}}^{\text{repeat}}$$

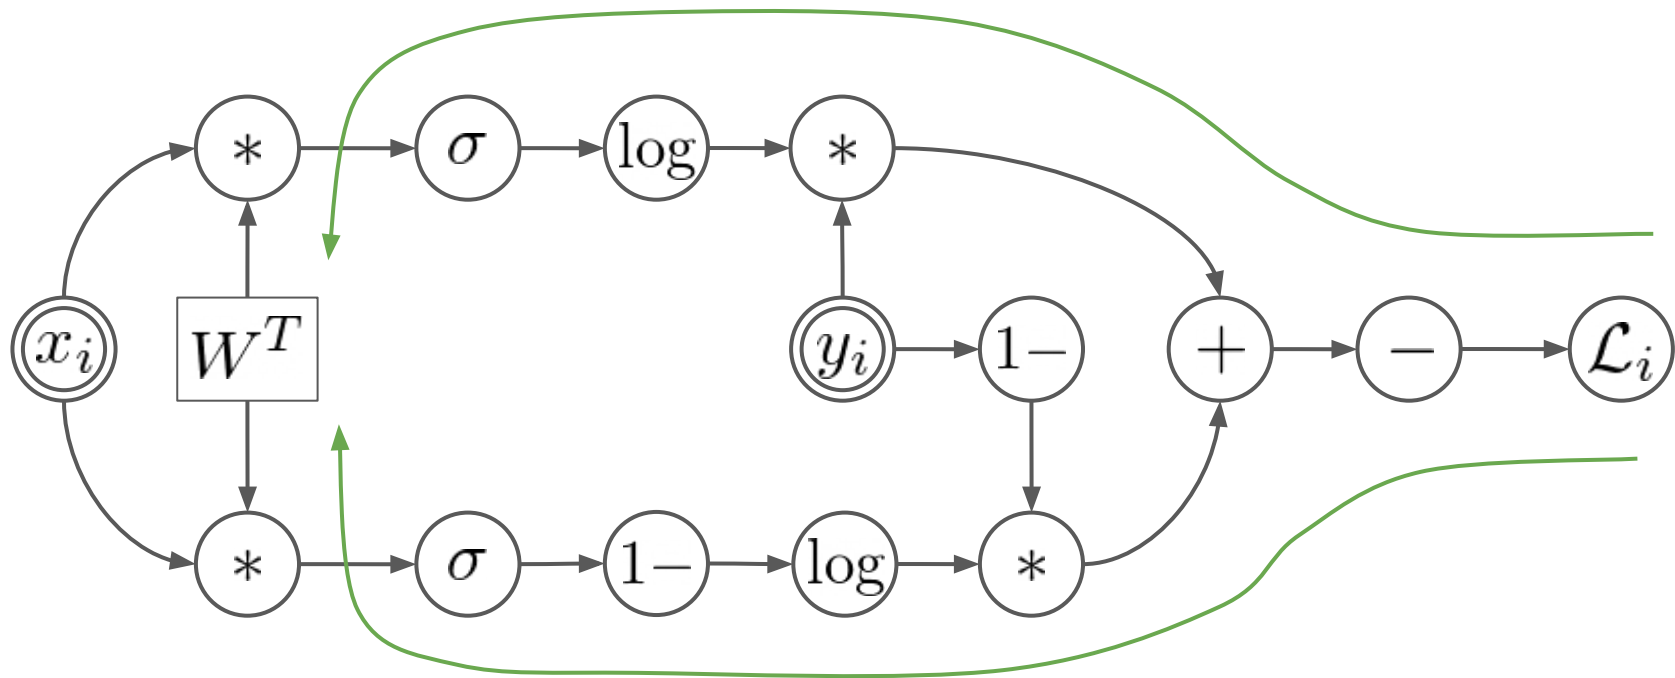


Computational Graph

$$\frac{\partial \mathcal{L}_i}{\partial W^T} = \frac{\partial \mathcal{L}_i}{\partial z} \left(\frac{\partial z}{\partial u} \frac{\partial u}{\partial W^T} + \frac{\partial z}{\partial v} \frac{\partial v}{\partial W^T} \right)$$




Computational Graph: Backprop



Computational Graph


Forward Pass: Compute output



Backward Pass: Compute gradients


Logistic Regression

$$p(y_i = 1|x_i, W) = \sigma(x_i W^T)$$

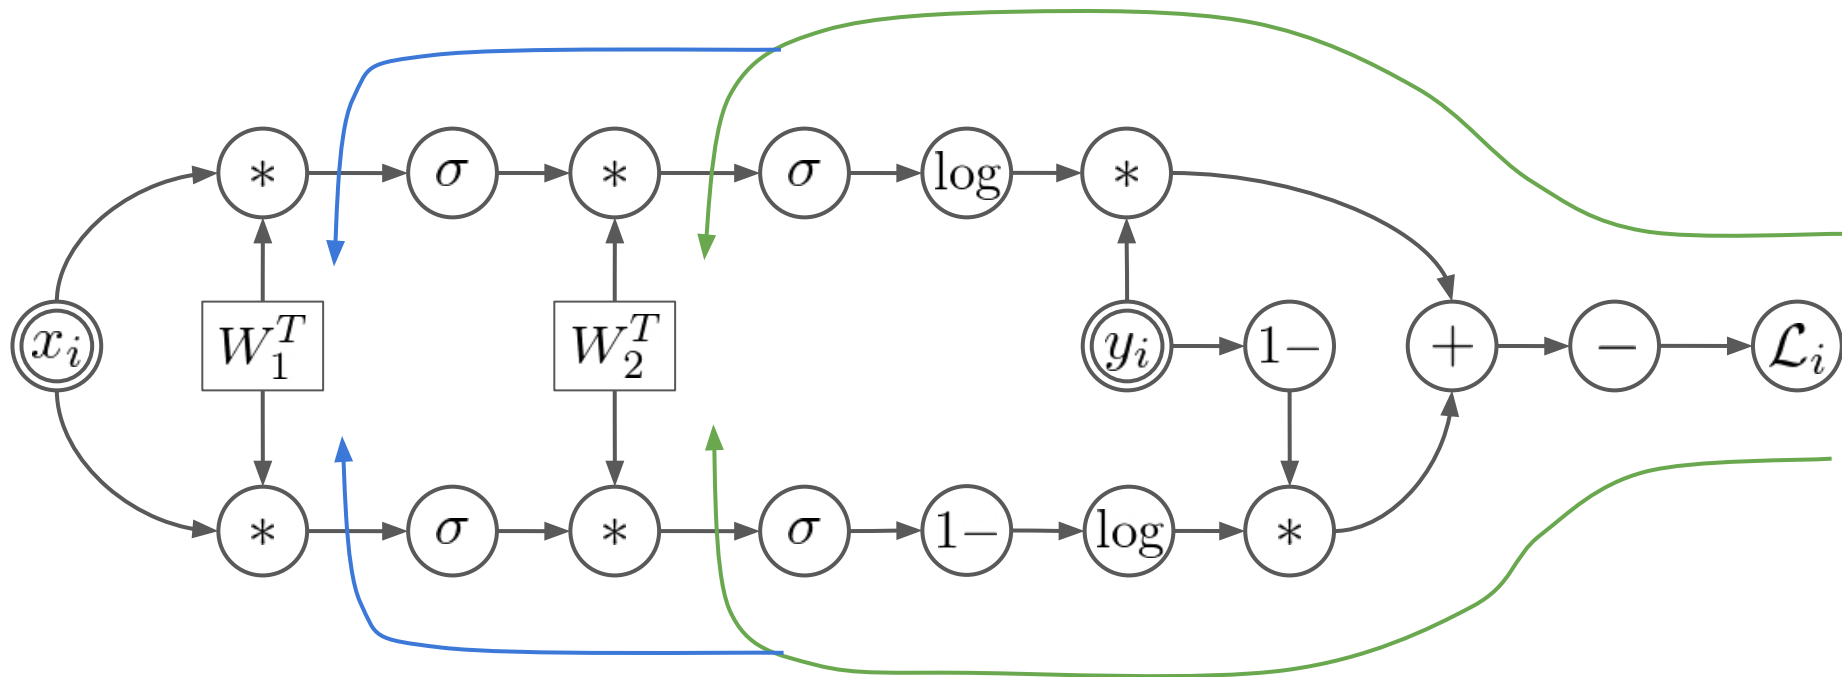


2-layered Logistic Regression

$$p(y_i = 1|x_i, W_1, W_2) = \sigma(\sigma(x_i W_1^T) W_2^T)$$




2-layered Logistic Regression



1-layered vs 2-layered Logistic Regression

Demo...

Deep Neural Networks: ResNet50



Differentiable Programming

Automatic differentiation:

1. Write a function: `def f(...)`
2. Automatically generate: `grad_f(...)`

Parameter learning:

- lots of algorithms based on stochastic gradient descent

PYTORCH^{*}

Three main components:

- Tensor processing library both for GPU & CPU
- Autograd library
- Neural networks library

^{*} Many deep learning and autodiff frameworks exist, PyTorch is the one I currently use the most.

2-layer Logistic Regression with Numpy

```
W1 = np.random.randn(H, D + 1)
W2 = np.random.randn(1, H)

for iter in range(1000):
    nu = sigmoid(X @ W1.T)
    mu = sigmoid(nu @ W2.T)

    nll = - y.T @ np.log(mu) - (1 - y).T @ np.log(1 - mu)

    gW1 = X.T @ (((mu - y) @ W2) * nu * (1 - nu))
    W1 -= lr * gW1.T

    gW2 = nu.T @ (mu - y)
    W2 -= lr * gW2.T
```

2-layer Logistic Regression with Torch

```
W1 = torch.randn(H, D + 1, requires_grad=True)
W2 = torch.randn(1, H, requires_grad=True)

for iter in range(1000):
    nu = torch.sigmoid(X @ W1.T)
    mu = torch.sigmoid(nu @ W2.T)
    nll = - y.T @ np.log(mu) - (1 - y).T @ np.log(1 - mu)
    nll.squeeze().backward()

    with torch.no_grad():
        W1 -= lr * W1.grad
        W2 -= lr * W2.grad
```

2-layer Logistic Regression with Torch.nn

```
class LogisticRegression(nn.Module):  
    def __init__(self, input_dims, hidden_dims):  
        super(LogisticRegression, self).__init__()  
        self.lin1 = nn.Linear(input_dims, hidden_dims)  
        self.lin2 = nn.Linear(hidden_dims, 1)  
  
    def forward(self, x):  
        h = F.sigmoid(self.lin1(x))  
        y = F.sigmoid(self.lin2(h))  
        return y
```

2-layer Logistic Regression with Torch.nn

```
model = LogisticRegression(D, H)
optimizer = torch.optim.Adadelta(model.parameters(), lr=lr)

for iter in range(1000):
    nll = F.nll(model(X))
    nll.backward()
    optimizer.step()
    optimizer.zero_grad()
```


3-layer Logistic Regression with Torch.nn

```
class LogisticRegression(nn.Module):
    def __init__(self, input_dims, hidden_dims):
        super(LogisticRegression, self).__init__()
        self.lin1 = nn.Linear(input_dims, hidden_dims)
        self.lin2 = nn.Linear(hidden_dims, hidden_dims)
        self.lin3 = nn.Linear(hidden_dims, 1)

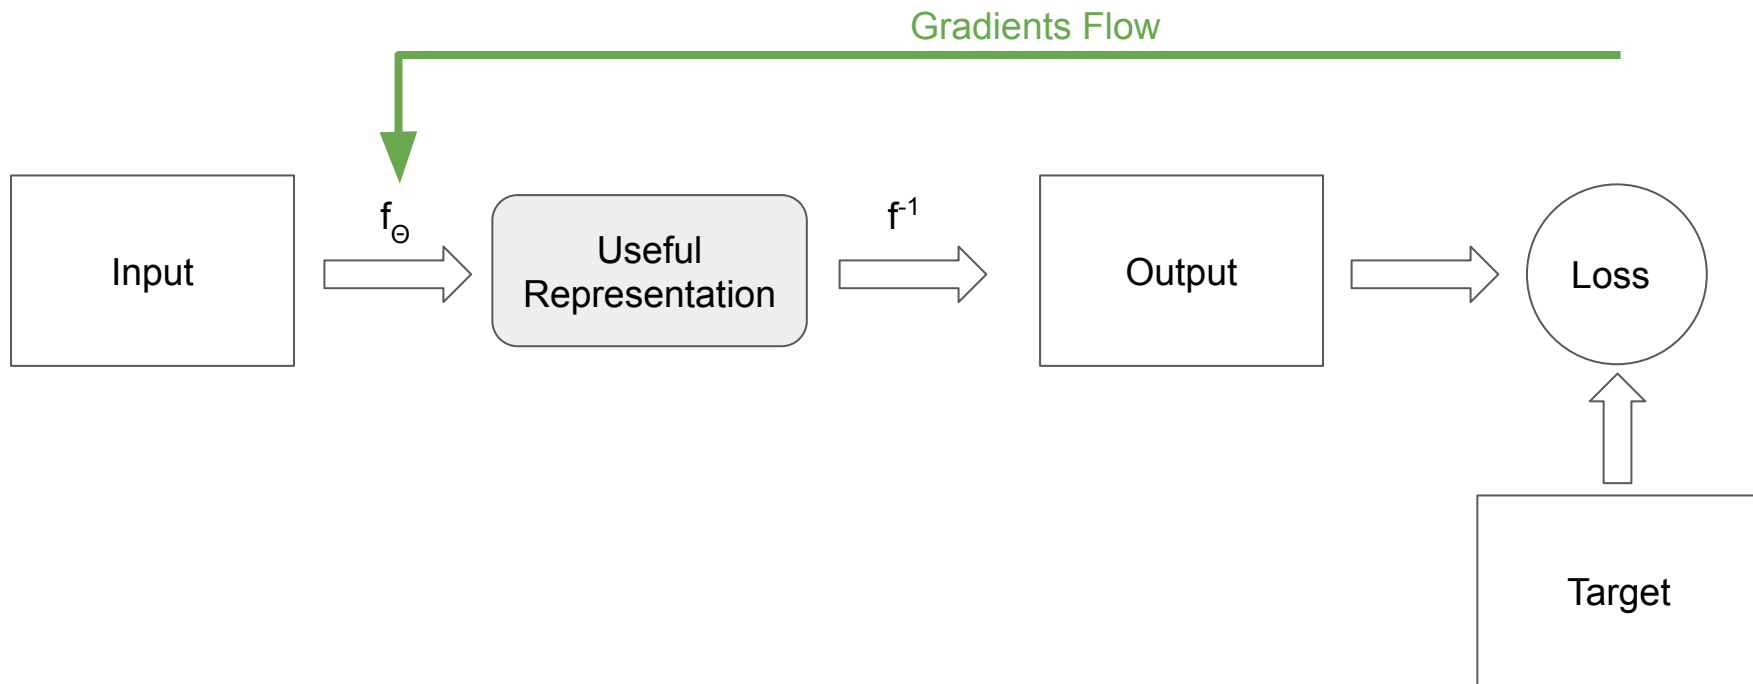
    def forward(self, x):
        h = F.sigmoid(self.lin1(x))
        h = F.sigmoid(self.lin2(h))
        y = F.sigmoid(self.lin3(h))
        return y
```

N-layer Logistic Regression with Torch.nn

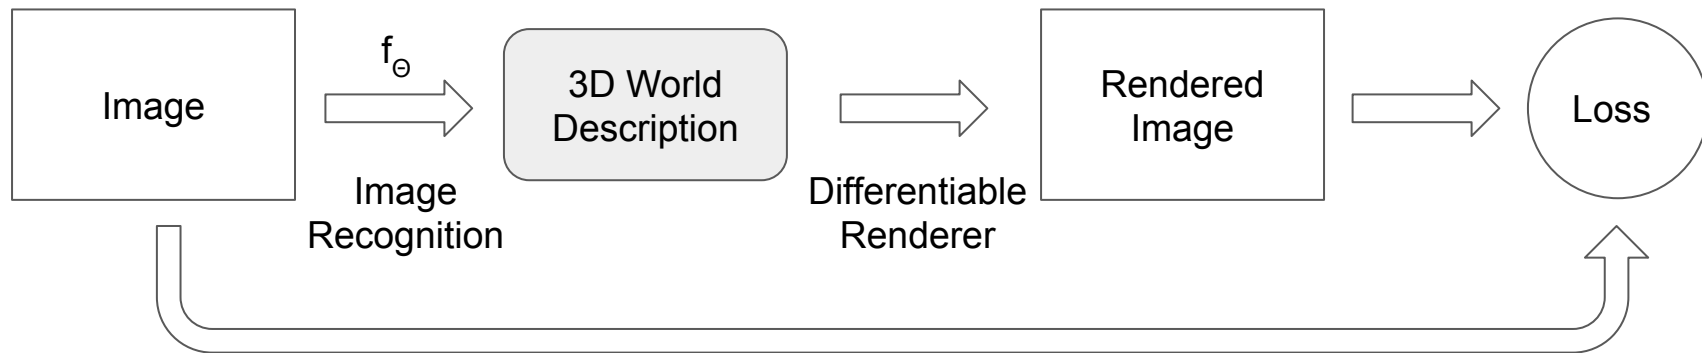
```
class LogisticRegression(nn.Module):
    def __init__(self, dims: List[int]):
        super(LogisticRegression, self).__init__()
        self.lins = [
            nn.Linear(in_dims, out_dims)
            for in_dims, out_dims in zip(dims[:-1], dims[1:])
        ]

    def forward(self, x):
        res = x
        for lin in self.lins:
            res = lin(res)
        return res
```

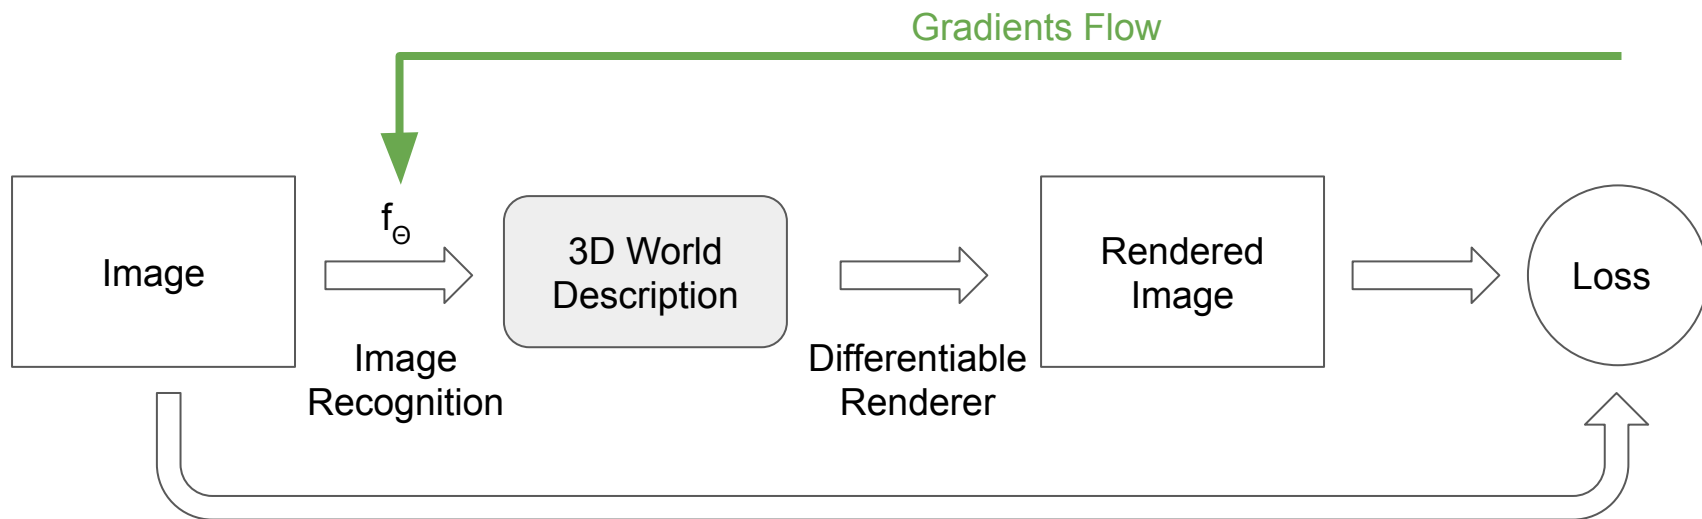
Learning Inverse Functions



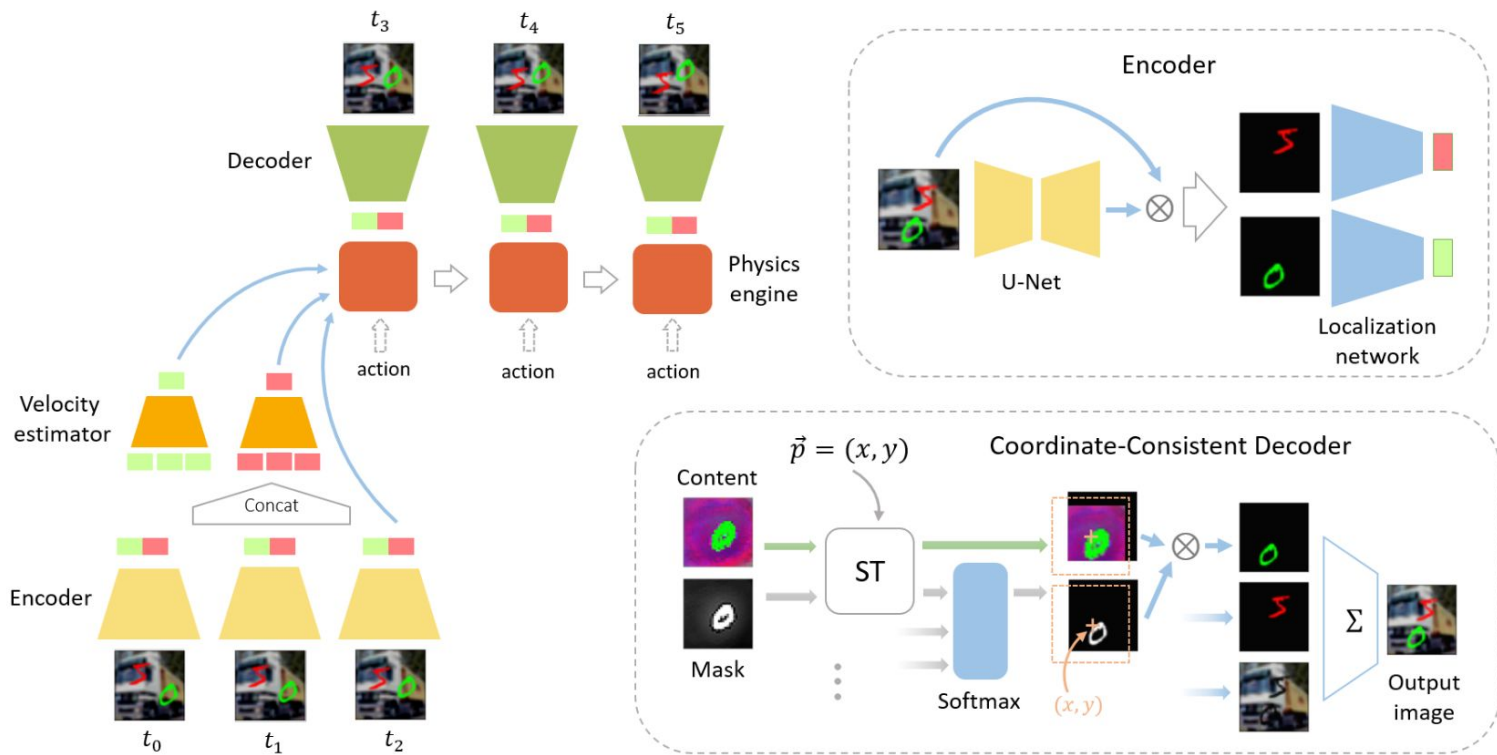
Differentiable Renderer (Inverse Graphics)




Differentiable Renderer (Inverse Graphics)




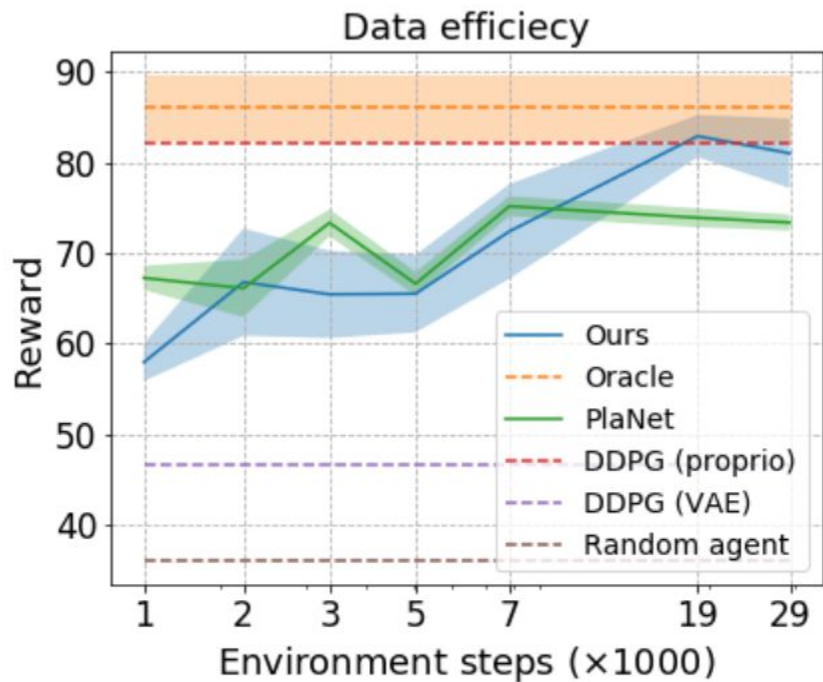
Physics-as-Inverse-Graphics



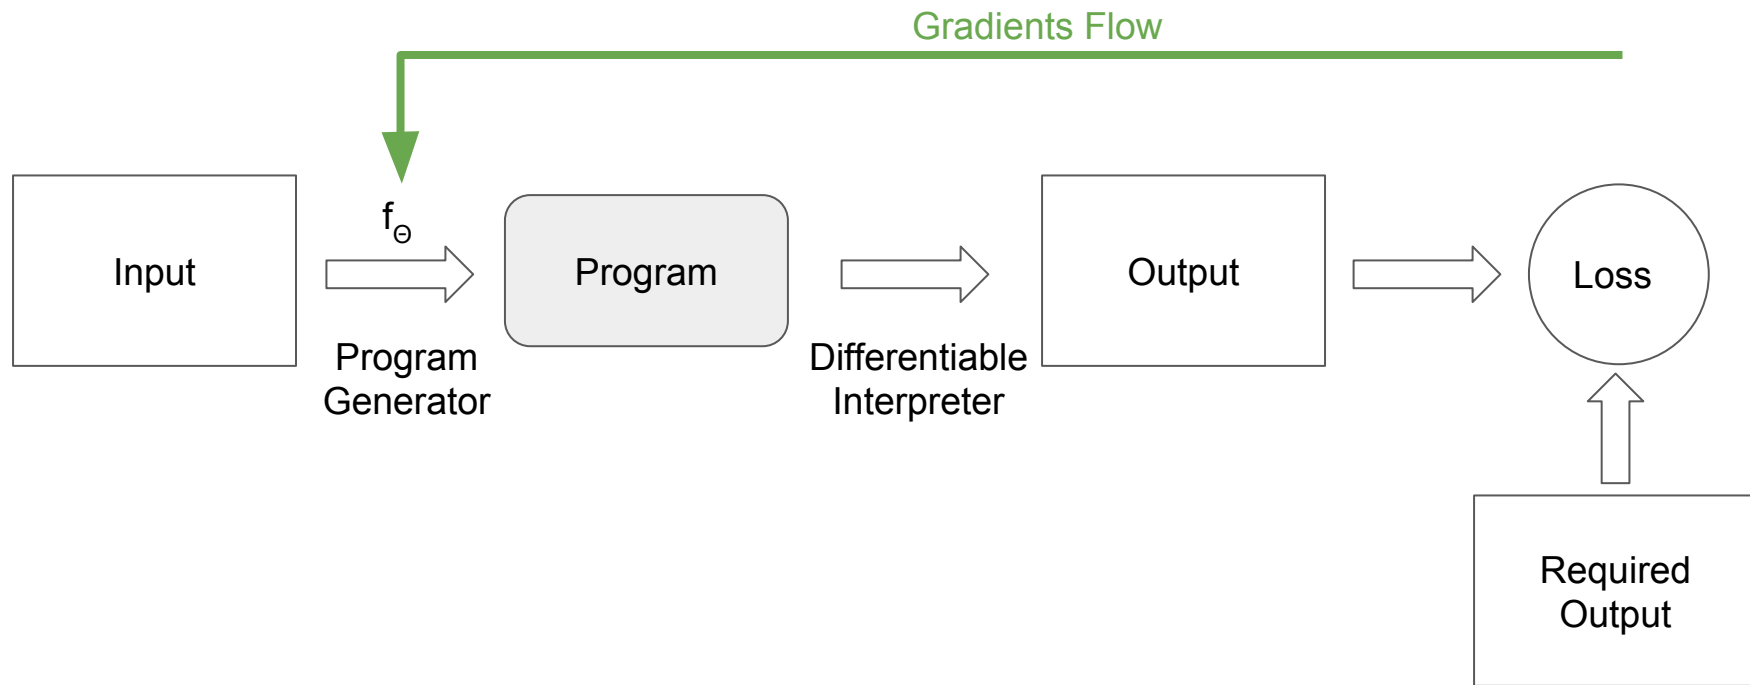
Learn Physical Parameters



Vision Based Model Predictive Control



Differentiable Program Interpreter

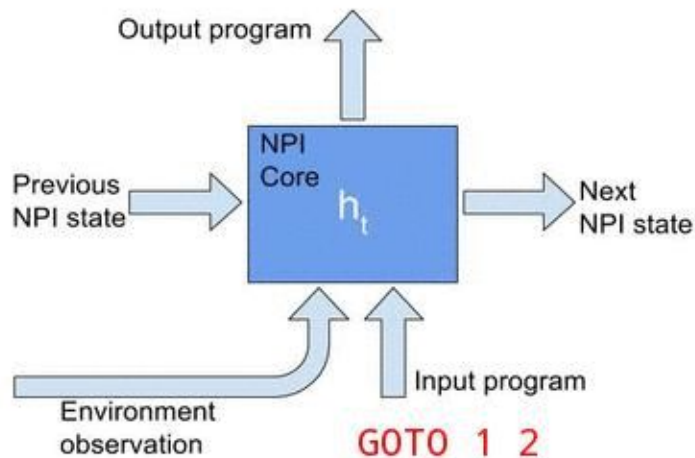


Neural Program Interpreter

Car rendering



NPI inference



Generated commands

GOTO 1 2

Differentiable Program with Neural Libraries

Declaration & initialization

```
# constants
max_int = 15; n_instr = 3; T = 45
W = 5; H = 3; w = 50; h = 50

# variables
img_grid = InputTensor(w, h)[W, H]
init_X = Input(W)
init_Y = Input(H)
final_X = Output(W)
final_Y = Output(Y)
path_len = Output(max_int)

instr = Param(4)[n_instr]
goto = Param(n_instr)[n_instr]

X = Var(W)[T]
Y = Var(H)[T]
dir = Var(5)[T]
reg = Var(max_int)[T]
instr_ptr = Var(n_instr)[T]

X[0].set_to(init_X)
Y[0].set_to(init_Y)
dir[0].set_to(1)
reg[0].set_to(0)
instr_ptr[0].set_to(0)
```

Instruction Set

```
# Discrete operations
@Runtime([max_int], max_int)
def INC(a):
    return (a + 1) % max_int

@Runtime([max_int], max_int)
def DEC(a):
    return (a - 1) % max_int

@Runtime([W, 5], W)
def MOVE_X(x, dir):
    if dir == 1: return (x + 1) % W # →
    elif dir == 3: return (x - 1) % W # ←
    else: return x

@Runtime([H, 5], H)
def MOVE_Y(y, dir):
    if dir == 2: return (y - 1) % H # ↑
    elif dir == 4: return (y + 1) % H # ↓
    else: return y

# Learned operations
@Learn([Tensor(w, h)], 5,
        hid_sizes=[256, 256])
def LOOK(img):
    pass
```

Execution model

```

for t in range(T - 1):
    if dir[t] == 0: # halted
        if dir[t + 1].set_to(dir[t])
            X[t + 1].set_to(X[t])
            Y[t + 1].set_to(Y[t])
            reg[t + 1].set_to(reg[t])
    else:
        with instr_ptr[t] as i:
            if instr[i] == 0: # INC
                reg[t + 1].set_to(inc(reg[t + 1]))
            if instr[i] == 1: # DEC
                reg[t + 1].set_to(dec(reg[t + 1]))
            else:
                reg[t + 1].set_to(reg[t])

            if instr[i] == 2: # MOVE
                X[t + 1].set_to(MOVE_X(X[t], dir[t]))
                Y[t + 1].set_to(MOVE_Y(Y[t], dir[t]))
            else:
                X[t + 1].set_to(X[t])
                X[t + 1].set_to(Y[t])

            if instr[i] == 3: # LOOK
                with pos[t] as p:
                    dir[t + 1].set_to(LOOK(img_grid[p]))
            else:
                dir[t + 1].set_to(dir[t])


    instr_ptr[t + 1].set_to(goto[i])

final_X.set_to(X[T - 1])
final_Y.set_to(Y[T - 1])
path_len.set_to(reg)

```

Solution

```
instr = [3,2,0]
goto = [1,2,0]
```



Input-output
data set



```
init_X = 0
init_Y = 1

final_X =
final_Y =
path len =
```

Differentiable Neural Computer

